

Introduction to Parallel Processing

Isaac van Til
Texas State University - CS 3339
Fall 2018



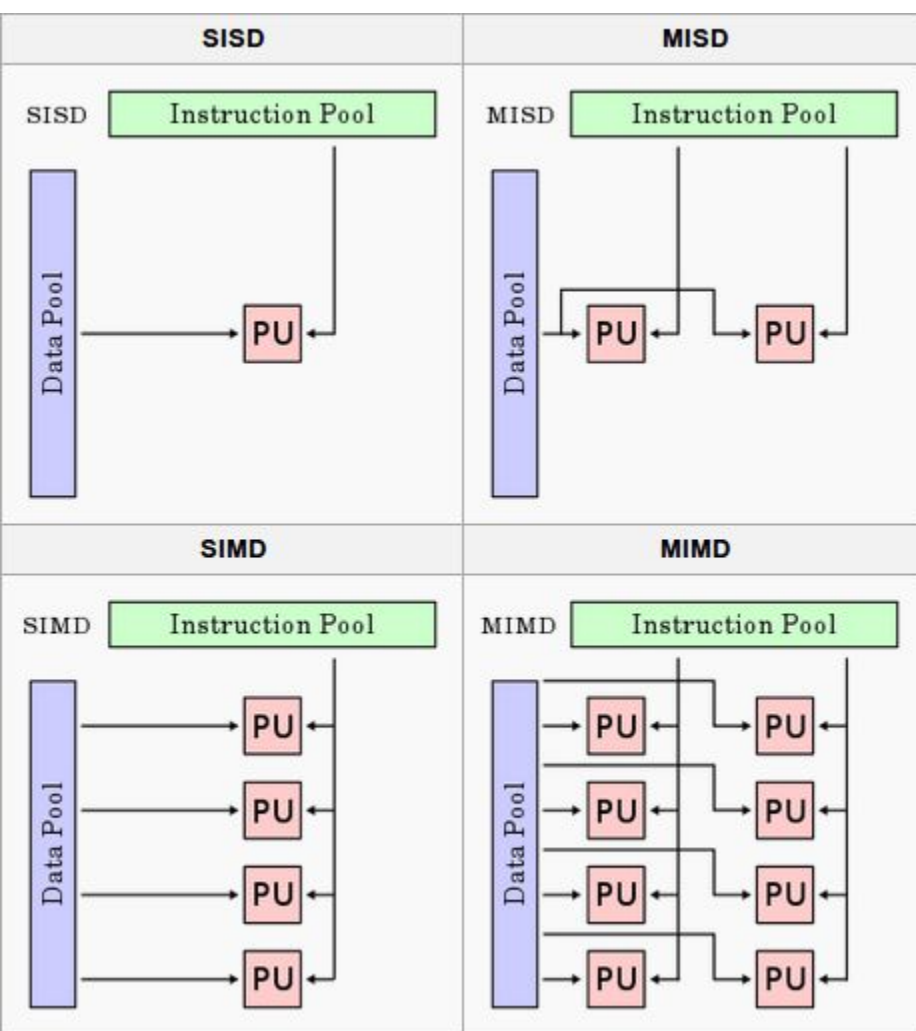
Overview

- Superpipeline / Superscalar
- Multi-core architecture
- Introduction to OpenMP and Multithreading



1. Superscalar Processors



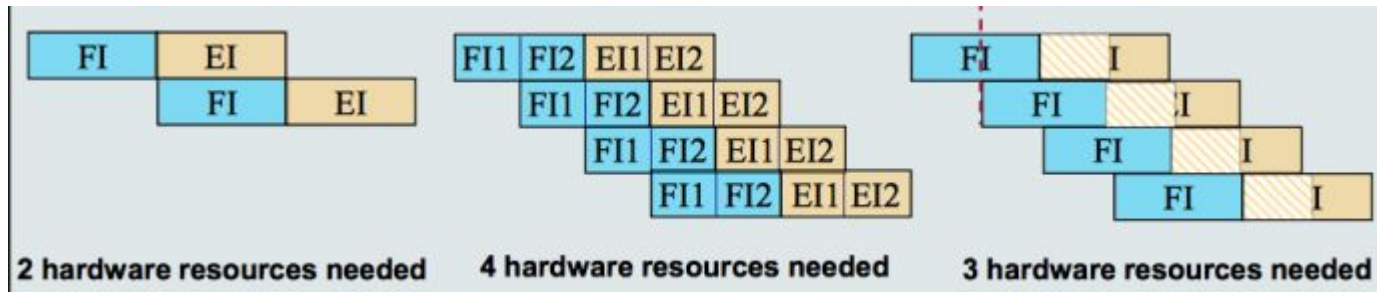


Flynn's Taxonomy



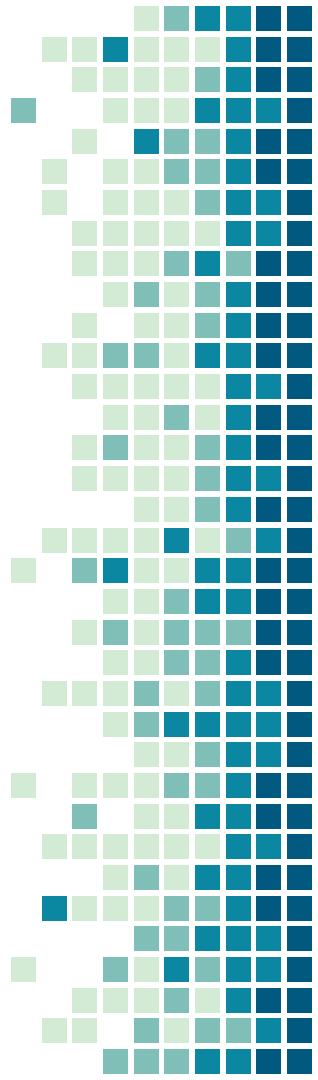
Superpipelining

- Divide the stages of a pipeline into several sub-stages, increases the number of instructions which are handled by the pipeline at the same time



Superpipelining Issues

- Increasing the number of stages/substages beyond an optimal limit reduces performance
- Not all stages can be divided into equal-length sub-stages
- Much more complex hardware
- Hazards more difficult to resolve
- Clock skew problem



The “Flynn Bottleneck”

- Single issue performance limit is $CPI = 1$
- With hazards, $CPI \geq 1$
- Diminishing returns from superpipelining




Solution:

- Issue multiple instructions per cycle
- Instruction-level parallelism (ILP)
- Superscalar architecture

	1	2	3	4	5	6	7
inst0	F	D	X	M	W		
inst1	F	D	X	M	W		
inst2		F	D	X	M	W	
inst3		F	D	X	M	W	

Superscalar vs. Superpipeline

- Base machine: 4-stage pipeline
 - Instruction fetch
 - Operation decode
 - Operation execution
 - Result write back
- Superpipeline of degree 2
 - A sub-stage often takes half a clock cycle to finish.
- Superscalar of degree 2
 - Two instructions are executed concurrently in each pipeline stage.
 - Duplication of hardware is required by definition.



Instruction Level Parallelism

- The classic von Neumann processors operate through **control flow** (instructions following each other linearly without regard for what data they involve)
- Superscalar processors utilize **data flow**, analyzing several instructions to find data dependencies and executing instructions in parallel that do not depend on each other. This is Instruction Level Parallelism.



Instruction Level Parallelism Mechanisms

- Multiple-issue
- Pipelining
- Branch prediction
- Out-of-order execution
- Prefetching



Introduction to Superscalar Architecture (SSA)

- SSA improves the performance of the execution of scalar instructions
- Several scalar instructions can be initiated simultaneously and executed independently
- Includes all the features in pipelining, but can execute multiple instructions in each pipeline stage



Introduction to Superscalar Architecture (SSA)

- Processors studied so far are fundamentally limited to $CPI \geq 1$
- Superscalar processors enable $CPI < 1$ ($IPC > 1$) by executing multiple instructions in parallel
- Can have both in-order and out-of-order superscalar processors

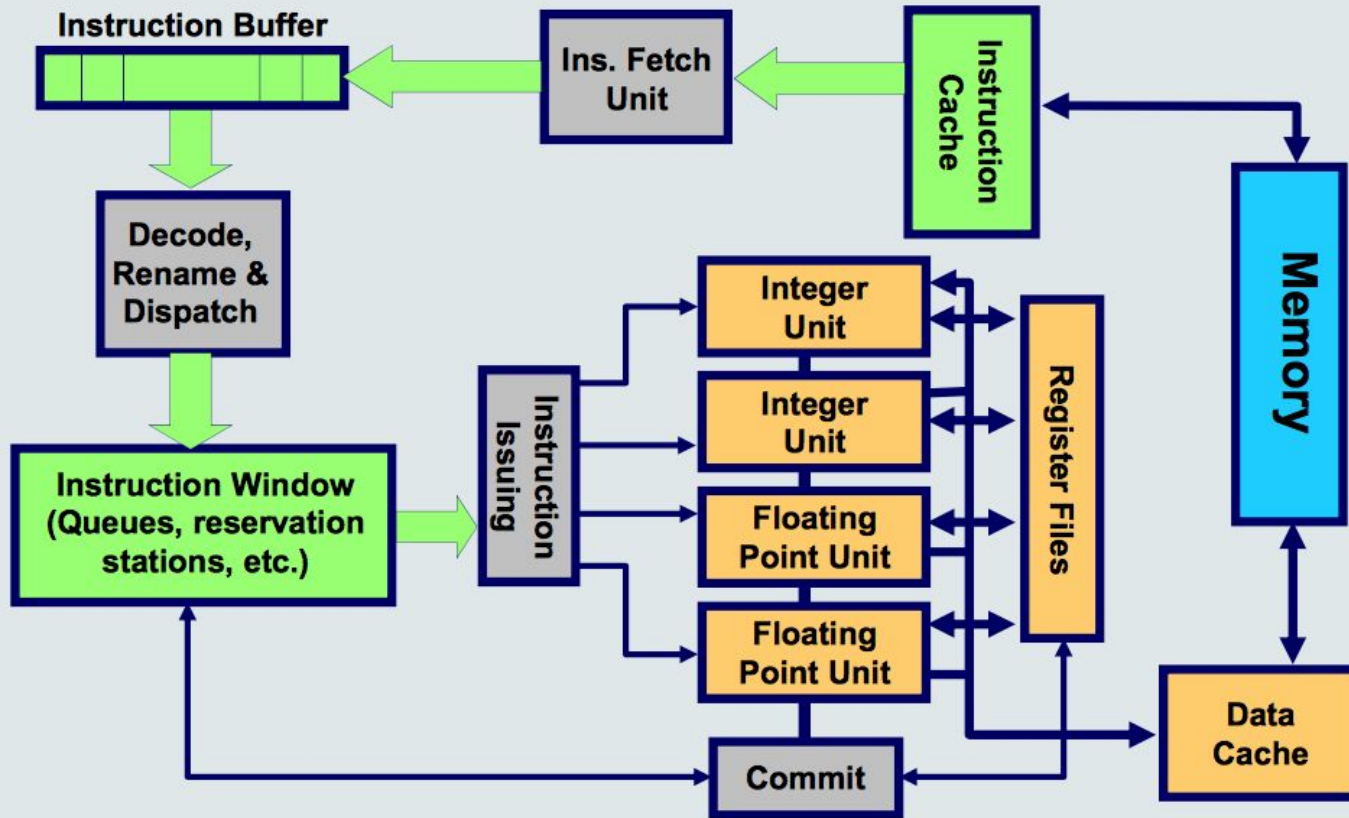


Implementation (in-order)

- Several instructions are issued and completed per clock cycle
- Several pipelines working in parallel
- Depending on the number and kind of parallel units available, a certain number of instructions can be executed in parallel
- Each unit is also pipelined and can execute several operations in different pipeline stages



An SSA Example



Implementation (out-of-order)

- A SSA processor fetches multiple instructions at a time, and attempts to find nearby instructions that are independent of each other
- Based on dependency analysis, the processor may issue and execute instructions in an order that differs from the original machine code
- The processor may eliminate some unnecessary dependencies by the use of additional registers and renaming of register references



Advantages of Superscalar

- Hardware solves everything (detects possible parallelism, register renaming)
- Binary compatibility
 - Programs don't need to be changed if functional units are added (without changing the instruction set)
 - Old programs benefit, the hardware simply issues the instructions more efficiently



Problems with Superscalar


- Complex hardware
- Power consumption
- Limited capacity to detect large numbers of parallel instructions
- Limited degree of intrinsic parallelism (instructions requiring the same CPU resources)



Data Hazards - review

Data-dependence


$r_3 \leftarrow r_1 \text{ op } r_2$
 $r_5 \leftarrow r_3 \text{ op } r_4$



Read-after-Write
(RAW) hazard

Anti-dependence

$r_3 \leftarrow r_1 \text{ op } r_2$
 $r_1 \leftarrow r_4 \text{ op } r_5$



Write-after-Read
(WAR) hazard

Output-dependence

$r_3 \leftarrow r_1 \text{ op } r_2$
 $r_3 \leftarrow r_6 \text{ op } r_7$



Write-after-Write
(WAW) hazard

Superscalar Pipeline Diagrams - Ideal

scalar

	1	2	3	4	5	6	7	8	9	10	11	12
lw 0(r1)→r2	F	D	X	M	W							
lw 4(r1)→r3		F	D	X	M	W						
lw 8(r1)→r4			F	D	X	M	W					
add r14,r15→r6				F	D	X	M	W				
add r12,r13→r7					F	D	X	M	W			
add r17,r16→r8						F	D	X	M	W		
lw 0(r18)→r9							F	D	X	M	W	

2-way superscalar

	1	2	3	4	5	6	7	8	9	10	11	12
lw 0(r1)→r2	F	D	X	M	W							
lw 4(r1)→r3	F	D	X	M	W							
lw 8(r1)→r4		F	D	X	M	W						
add r14,r15→r6		F	D	X	M	W						
add r12,r13→r7			F	D	X	M	W					
add r17,r16→r8			F	D	X	M	W					
lw 0(r18)→r9				F	D	X	M	W				

Superscalar Pipeline Diagrams - Realistic

scalar

```
lw 0(r1) → r2
lw 4(r1) → r3
lw 8(r1) → r4
add r4, r5 → r6
add r2, r3 → r7
add r7, r6 → r8
lw 4(r8) → r9
```

	1	2	3	4	5	6	7	8	9	10	11	12
lw 0(r1) → r2	F	D	X	M	W							
lw 4(r1) → r3		F	D	X	M	W						
lw 8(r1) → r4			F	D	X	M	W					
add r4, r5 → r6				F	d*	D	X	M	W			
add r2, r3 → r7						F	D	X	M	W		
add r7, r6 → r8							F	D	X	M	W	
lw 4(r8) → r9								F	D	X	M	W

2-way superscalar

```
lw 0(r1) → r2
lw 4(r1) → r3
lw 8(r1) → r4
add r4, r5 → r6
add r2, r3 → r7
add r7, r6 → r8
lw 4(r8) → r9
```

	1	2	3	4	5	6	7	8	9	10	11	12
lw 0(r1) → r2	F	D	X	M	W							
lw 4(r1) → r3	F	D	X	M	W							
lw 8(r1) → r4		F	D	X	M	W						
add r4, r5 → r6		F	d*	d*	D	X	M	W				
add r2, r3 → r7			F	d*	D	X	M	W				
add r7, r6 → r8					F	D	X	M	W			
lw 4(r8) → r9					F	d*	D	X	M	W		

Potential improvements

- Very Long Instruction Word Processor (VLIW)
- Explicit Parallelism
- Loop unrolling

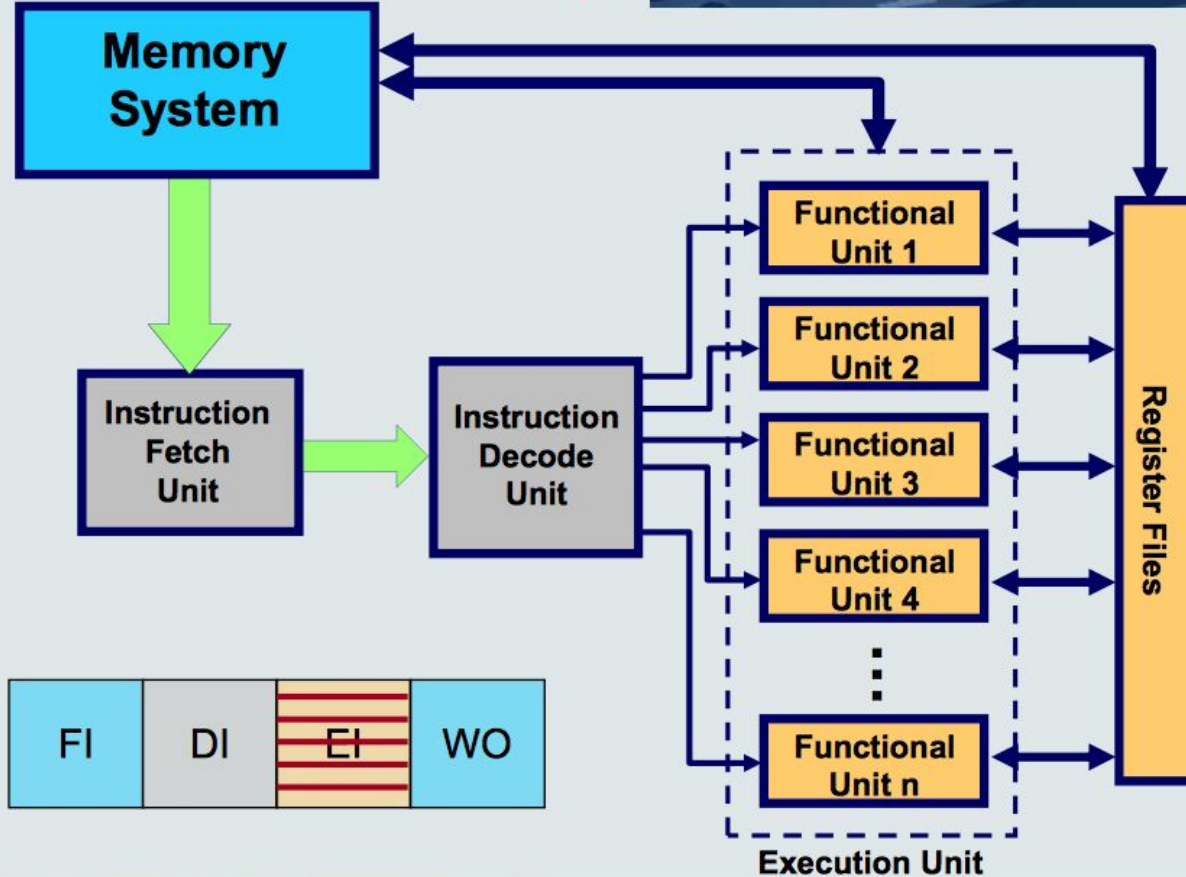


Very Long Instruction Word Processors

- Several operations that can be executed in parallel are placed in a single instruction word
- Compile-time parallelism detection
- After one instruction has been fetched all the corresponding operations are issued in parallel (no hardware needed for run-time detection of parallelism)
- Compiler can analyze the entire program to detect parallel operations



Typical Organization



Explicit Parallelism

- Instruction parallelism scheduled at compile time.
 - Included within the machine instructions explicitly.
- Processor uses this information to perform parallel execution.
- The hardware is less complex.
 - The controller is similar to a simple scalar computer.
 - The number of FUs can be increased without needing additional sophisticated hardware to detect parallelism, as in SSA.
- Compiler has much more time to determine possible parallel operations.
 - This analysis is only done once off-line, while run-time detection is carried out by SSA hardware for each execution of the code.
- Good compilers can detect parallelism based on global analysis of the whole program.



An Example

```
for (i=959; i>=0; i--)  
    x[i] = x[i] + s;
```

Assumptions:

x is an array of floating point values;
s is a floating point constant.

Memory allocation:

- R1 initially contains the address of the last element in **x**; the other elements are at lower addresses; **x**[0] is at address 0.
- Floating point register F2 contains the value **s**.
- A floating point value is 8 bytes long.

For an ordinary processor, this C code will be compiled to:

Loop:	LDD	F0, (R1)	F0 := x[i]; (load double)
	ADF	F4, F0, F2	F4 := F0 + F2; (add floating point)
	STD	(R1), F4	x[i] := F4; (store double)
	SBI	R1, R1, #8	R1 := R1 - 8;
	BGEZ	R1, Loop	branch if R1 ≥ 0.

An Example (Cont'd)

```

Loop: LDD    F0,(R1)    load double
      ADF    F4,F0,F2   add FP
      STD    (R1),F4    store double
      SBI    R1,R1,#8
      BGEZ   R1,Loop
    
```

Note the displacement of 8 for R1 is needed, because we have already subtracted 8 from R1

Cycl	Mem R	Mem R	FP 1	FP 2	I/BRA
1	LDD F0,(R1)				
2					
3			ADF F4,F0,F2		
4					
5					SBI R1,R1,#8
6	STD (R1+8),F4				BGEZ R1,Loop

- One iteration takes 6 cycles; the whole loop takes $960 \times 6 = 5760$ cycles.
- Almost no parallelism; most of the fields in the instructions are empty.
- There are two completely empty cycles.

Loop Unrolling

Let us rewrite the example:

```
for (i=959; i>=0; i-=2){  
    x[i] = x[i] + s;  
    x[i-1] = x[i-1] + s;  
}
```

Loop unrolling: a technique used in compilers in order to increase the potential of parallelism in a program.

— It supports more efficient code generation for processors with instruction level parallelism.

For an ordinary processor, this C new code will be compiled to:

Loop:	LDD	F0, (R1)	F0:=x[i]; (load double)
	ADF	F4, F0, F2	F4:=F0+F2; (add floating pnt)
	STD	(R1), F4	x[i]:=F4; (store double)
	LDD	F6, (R1-8)	F6:=x[i-1]; (load double)
	ADF	F8, F6, F2	F8:=F6+F2; (add floating pnt)
	STD	(R1-8), F8	x[i-1]:=F8; (store double)
	SBI	R1, R1, #16	R1:=R1-16;
	BGEZ	R1, Loop	branch if R1 ≥ 0.

Loop Unrolling (2 iterations)

Cycle

1	LDD F0,(R1)	LDD F6,(R1-8)			
2					
3			ADF F4,F0,F2	ADF F8,F6,F2	
4					
5					SBI R1,R1,#16
6	STD(R1+16),F4	STD(R1+8),F8			BGEZ R1,Loop

- There is an increased degree of parallelism in this case.
- We still have two completely empty cycles and many empty operations.
- However, we have basically double the performance:
 - Two iterations take 6 cycles
 - The whole loop takes $480 \times 6 = 2880$ cycles


2.

Multi-core architecture OpenMP & Multithreading



Moore's Law


transistors/chip doubles every 1.5 to 2 years



Consequences of Moore's Law

- Performance increased in tandem with transistor density
- People expect that performance and optimization comes from hardware
- Less emphasis is put on software optimization










Figure 1.28: Projected heat dissipation of a CPU if trends had continued – this graph courtesy Pat Helsinger

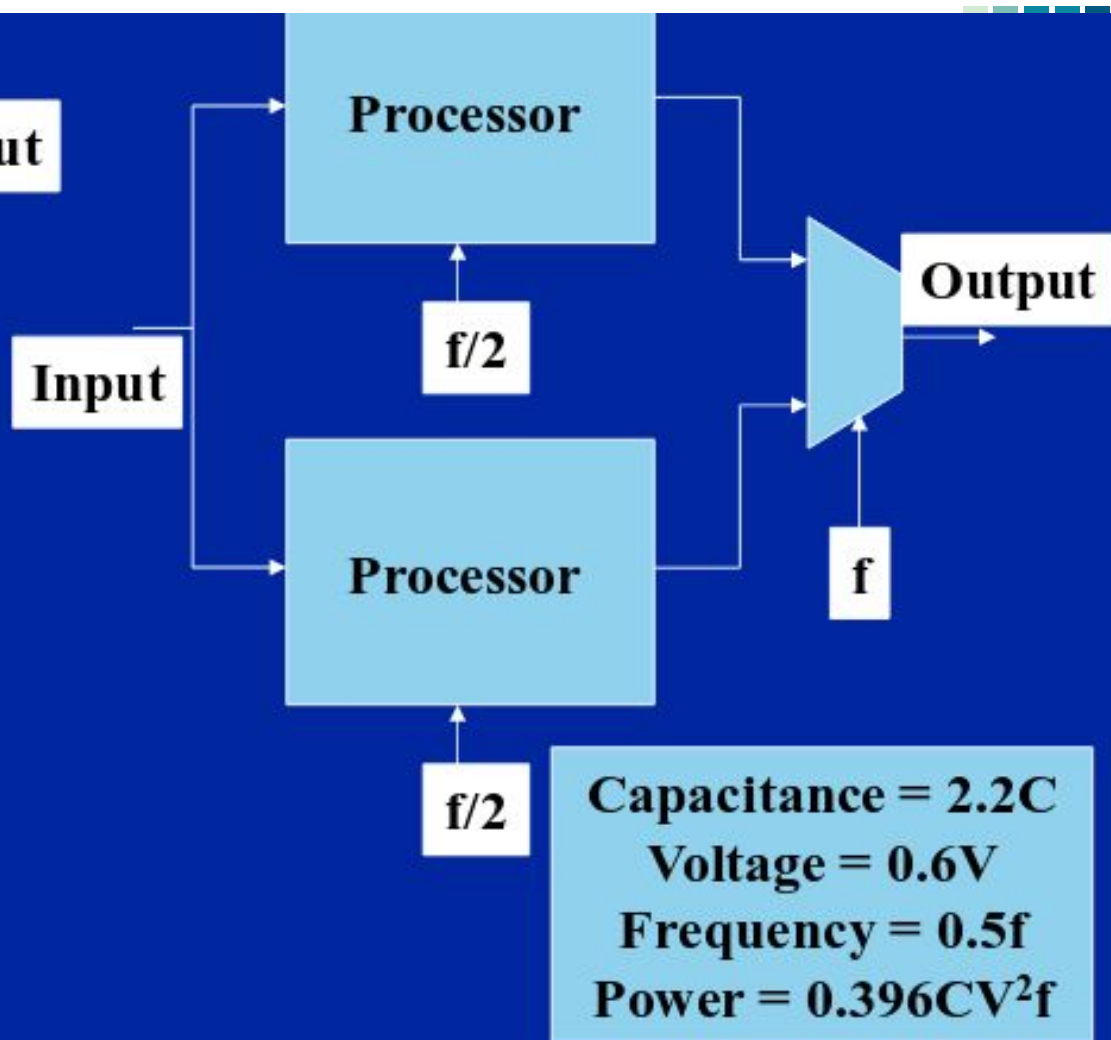
Solution:

- Create simpler cores designed around power optimization instead of performance optimization
- Use more than one of them!





Capacitance = C
Voltage = V
Frequency = f
Power = CV^2f



Capacitance = $2.2C$
Voltage = $0.6V$
Frequency = $0.5f$
Power = $0.396CV^2f$

Results:

- Performance now comes from software
- Parallel computing





966

Concurrency is when two or more tasks can start, run, and complete in overlapping time periods. It doesn't necessarily mean they'll ever both be running at the same instant. For example, *multitasking* on a single-core machine.




Parallelism is when tasks *literally* run at the same time, e.g., on a multicore processor.




Quoting [Sun's Multithreaded Programming Guide](#):

- Concurrency: A condition that exists when at least two threads are making progress. A more generalized form of parallelism that can include time-slicing as a form of virtual parallelism.
- Parallelism: A condition that arises when at least two threads are executing simultaneously.





Concurrent, non-parallel Execution



Concurrent, parallel Execution



An important distinction

- Concurrent problems are inherently concurrent (think web server) and you cannot even define the problem without concurrency
- Parallel problems have things happening at the same time, but you could do it without concurrency
- We will discuss **parallel** applications



Writing Parallel Software

- Step 1: find concurrency
- Step 2: organize into an algorithm that exploits parallelism
- Step 3: implement in a parallel language

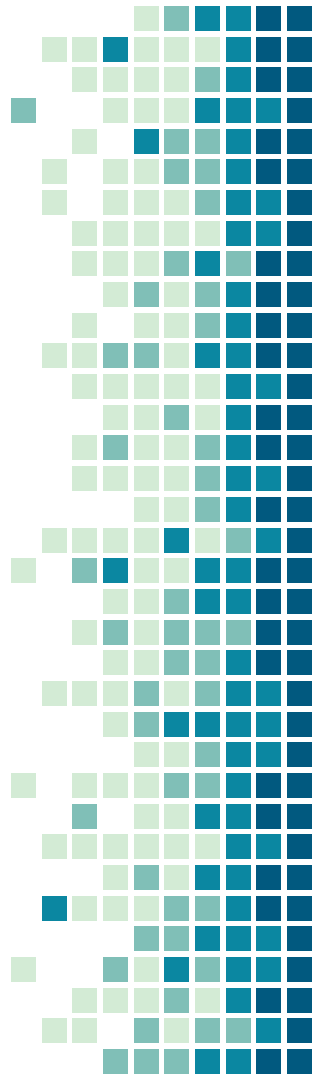


Example: Vector Addition

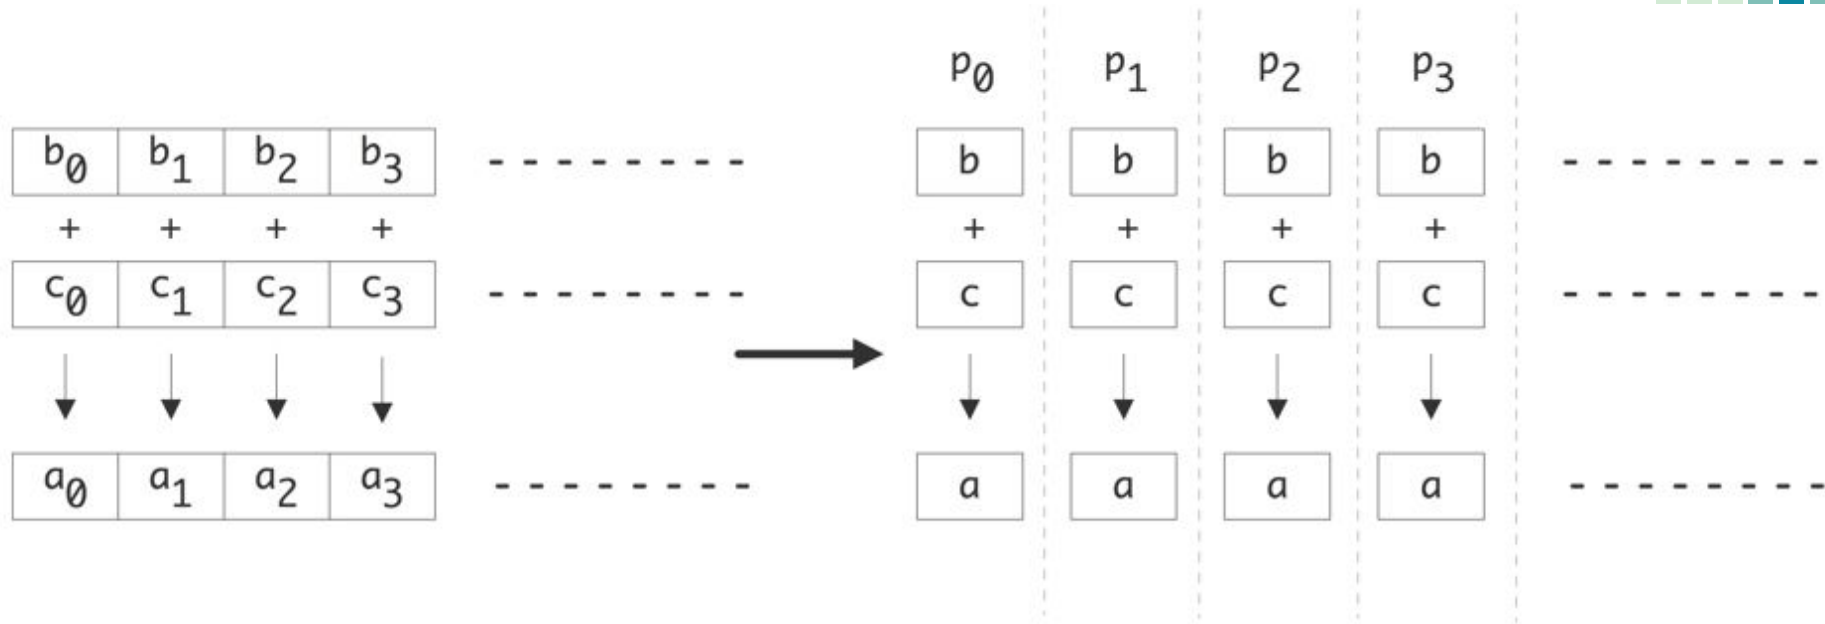
```
for (i=0; i<n; i++)  
    a[i] = b[i] + c[i];
```

Where is the concurrency?

How can we exploit parallelism?



Ideal case (n processors):

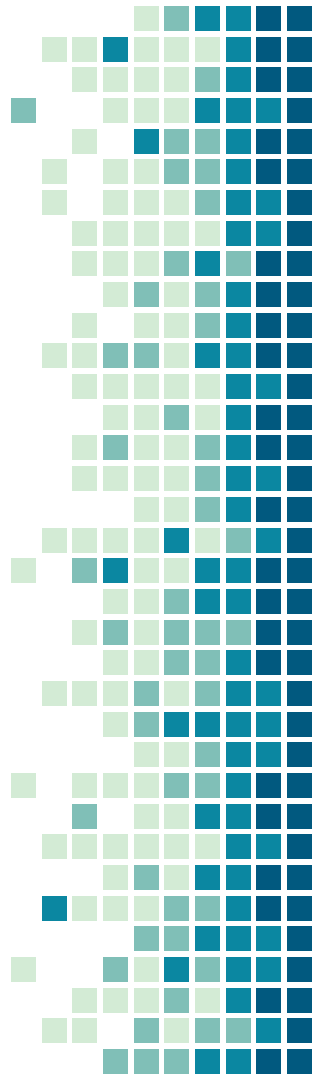


```
for (i=my_low; i<my_high; i++)  
    a[i] = b[i] + c[i];
```

The original algorithm takes n time.

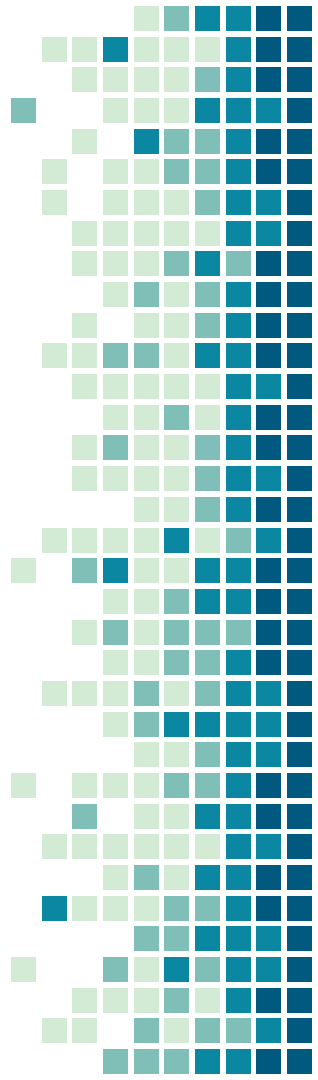
Parallel execution on p processors results in n/p

The parallel algorithm is faster by a factor of p .



This one is slightly less obvious

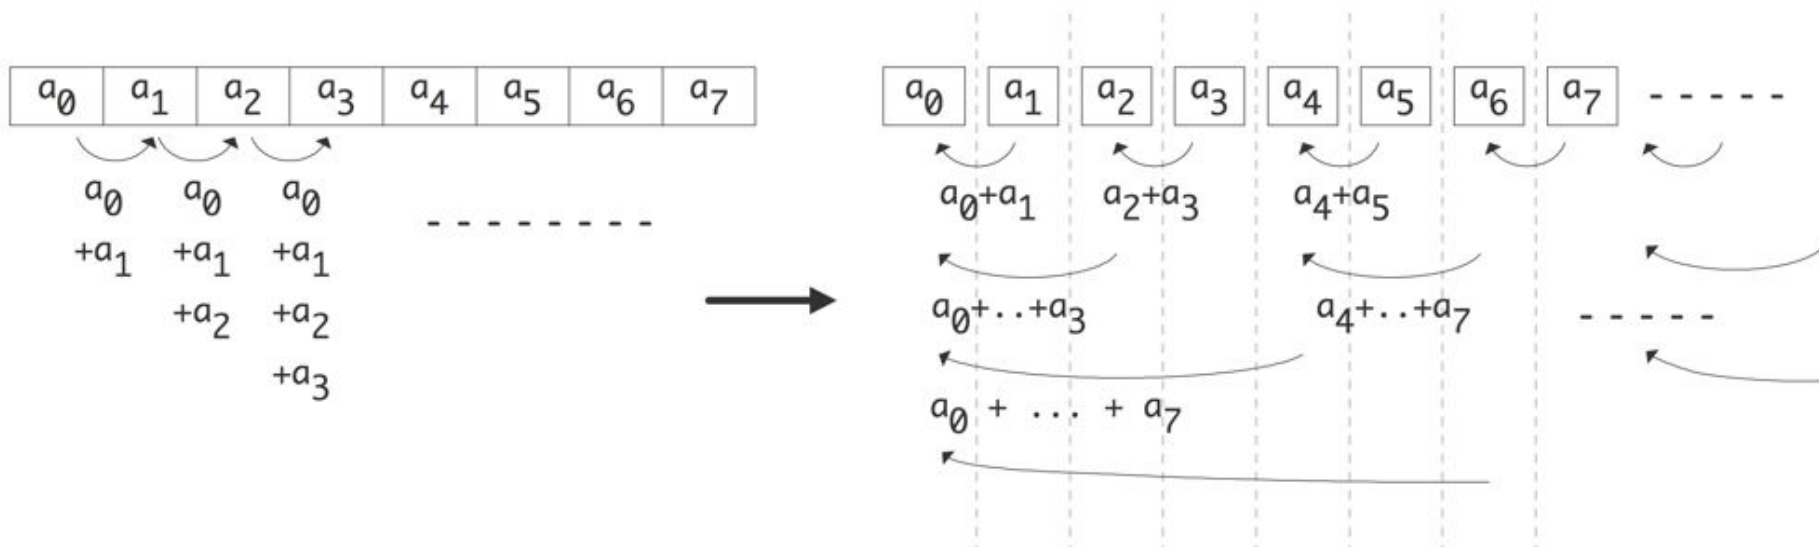

```
s = 0;  
for (i=0; i<n; i++)  
    s += x[i];
```


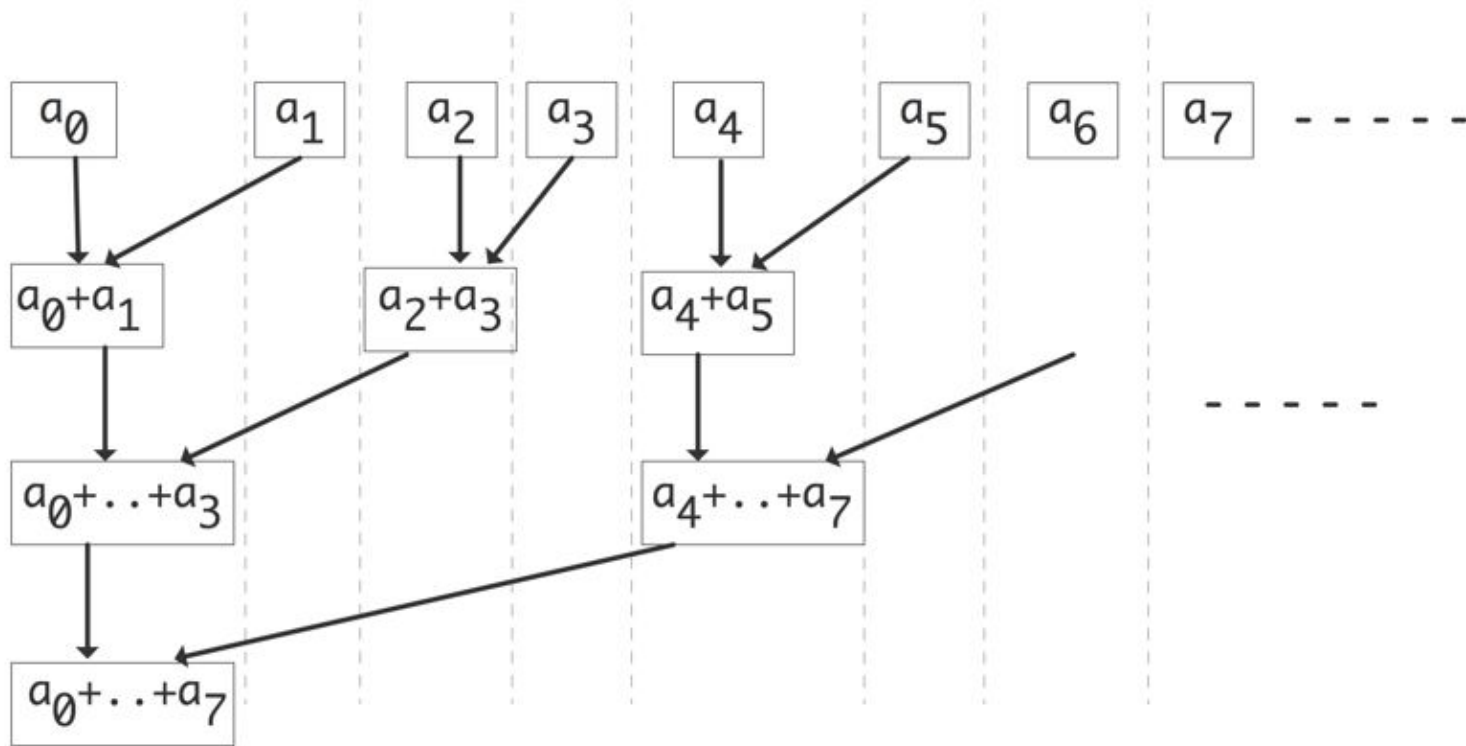


```

for (s=2; s<2*n; s*=2) {
    for (i=0; i<n-s/2; i+=s)
        x[i] += x[i+s/2];
}

```





Initial observations

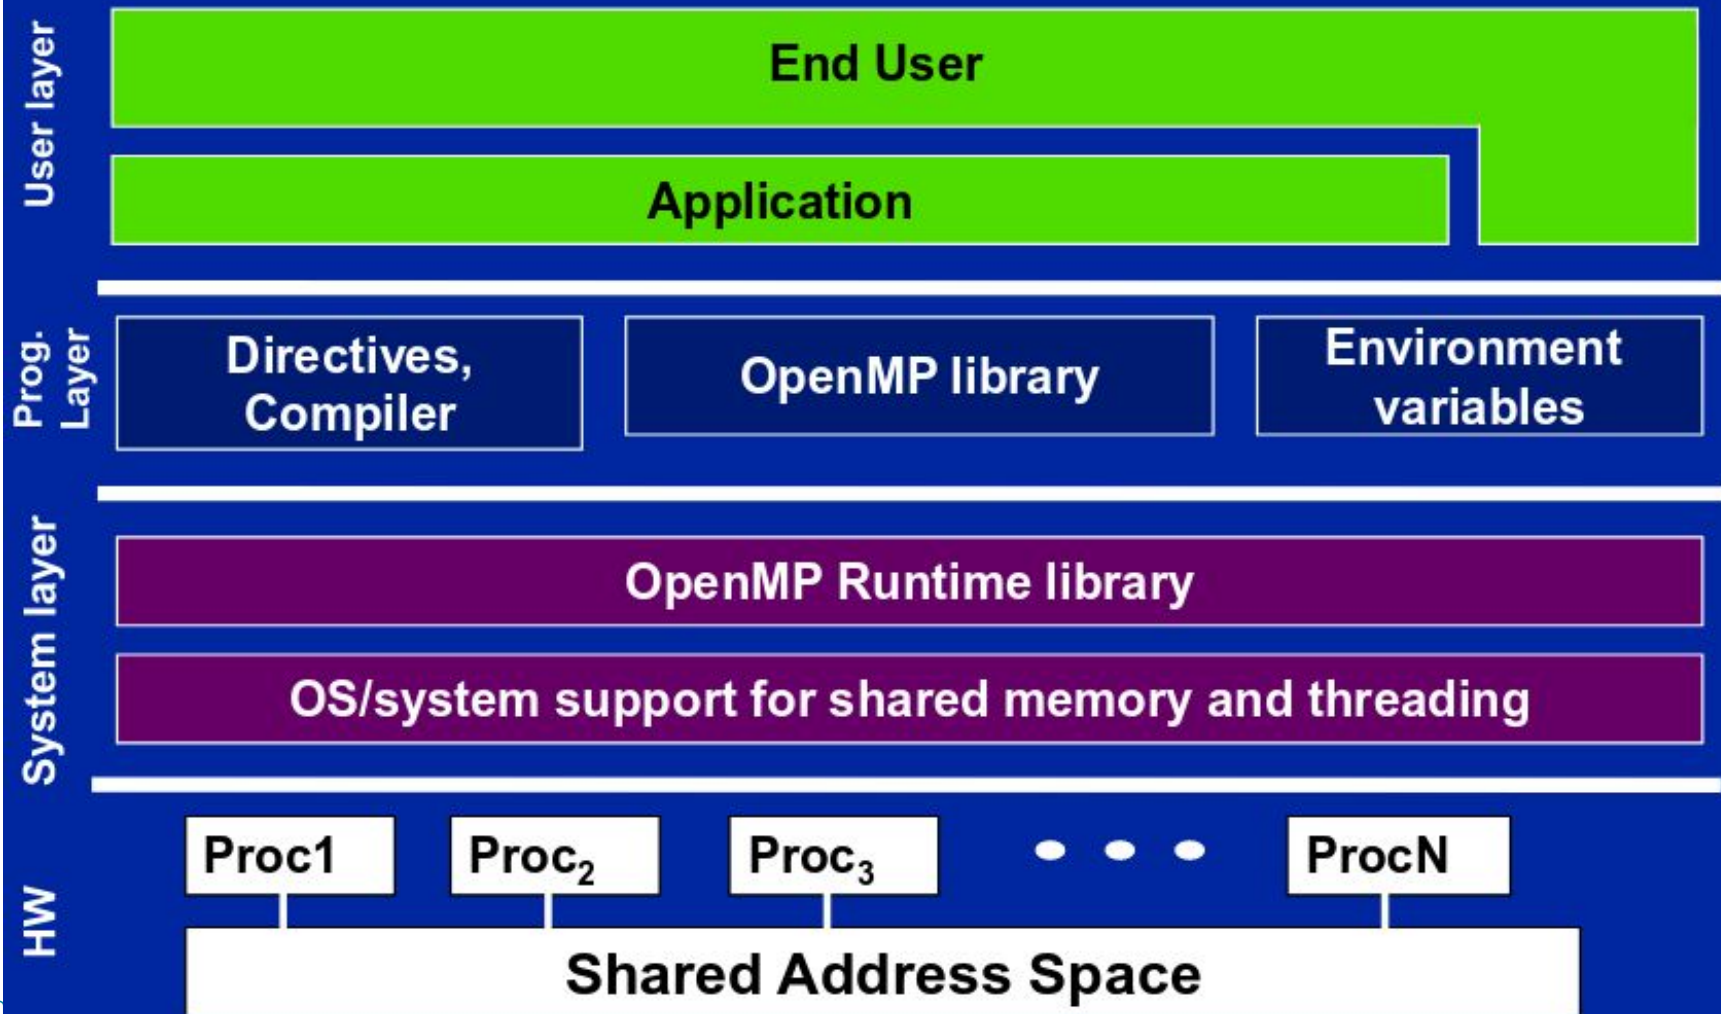
- Some algorithms need to be rewritten to make them parallel
- A parallel algorithm may not show perfect speedup
- Communicating data (example 2) between processors takes time



Open Multi-Processing (OpenMP)

- After you have a parallel algorithm, you can use a parallel language like OpenMP to implement it on a parallel computer
- OpenMP is a set of compiler directives and library routines for writing C/C++/FORTRAN in parallel
- Uses low level pthread stuff so you don't have to!





Basic Syntax

- `#include <omp.h>`
- **Most constructs are compiler directives**
 - `#pragma omp construct [clause[clause]...]`
- **Most constructs apply to a structured block**

```
#pragma omp parallel
{
    do_stuff();
}
```

```
#include <stdio.h>
#include <omp.h>


int main()
{
    #pragma omp parallel
    {
        int ID = omp_get_thread_num();
        printf("Hello from thread %d\n", ID);
    }

    return 0;
}
```

Compiling

```
$gcc -fopenmp hello.c
```

Other compilers- just Google it.



```
[isaac@X550ZA Parallel]$ ./a.out
Hello from thread 3
Hello from thread 0
Hello from thread 2
Hello from thread 1
[isaac@X550ZA Parallel]$ ./a.out
Hello from thread 3
Hello from thread 1
Hello from thread 2
Hello from thread 0
[isaac@X550ZA Parallel]$ ./a.out
Hello from thread 0
Hello from thread 1
Hello from thread 3
Hello from thread 2
[isaac@X550ZA Parallel]$ ./a.out
Hello from thread 2
Hello from thread 0
Hello from thread 3
Hello from thread 1
[isaac@X550ZA Parallel]$ ./a.out
Hello from thread 3
Hello from thread 0
Hello from thread 1
Hello from thread 2
[isaac@X550ZA Parallel]$ ./a.out
Hello from thread 3
Hello from thread 0
Hello from thread 2
Hello from thread 1
[isaac@X550ZA Parallel]$ ./a.out
Hello from thread 1
Hello from thread 0
Hello from thread 3
Hello from thread 2
```

Shared Memory Environment

- Symmetric Multiprocessing (SMP)
 - Shared address space with equal time access for each processor
 - Uniform Memory Access (UMA)
 - OS treats every processor the same way
- Non Uniform Memory Architecture (NUMA)
 - “Equal Time”
 - Cache architecture



Stack

funcA() var1
var2

Stack Pointer
Program Counter
Registers

Process

- An instance of a program execution.
- The execution context of a running program ... i.e. the resources associated with a program's execution.

text

main()
funcA()
funcB()
.....


data

array1
array2

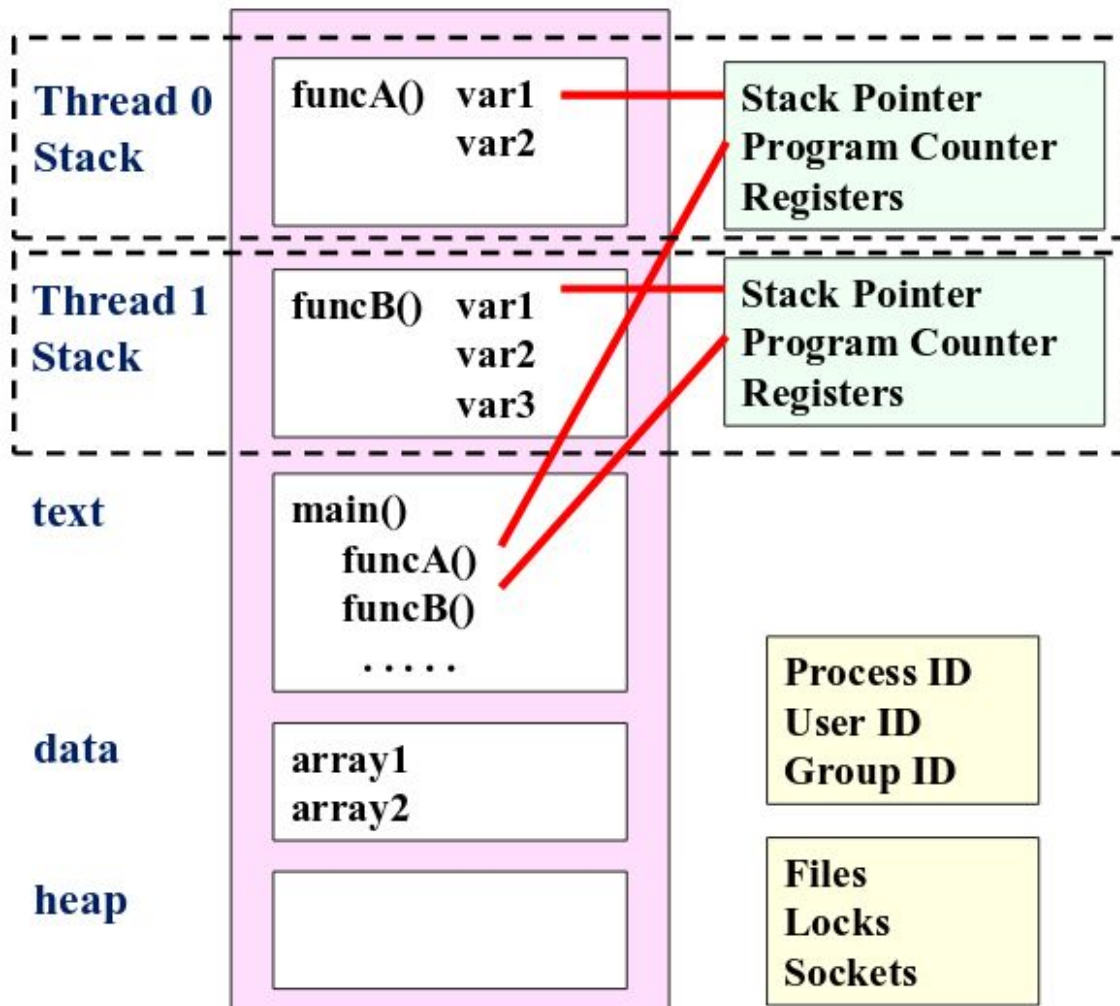
Process ID
User ID
Group ID

heap

Files
Locks
Sockets





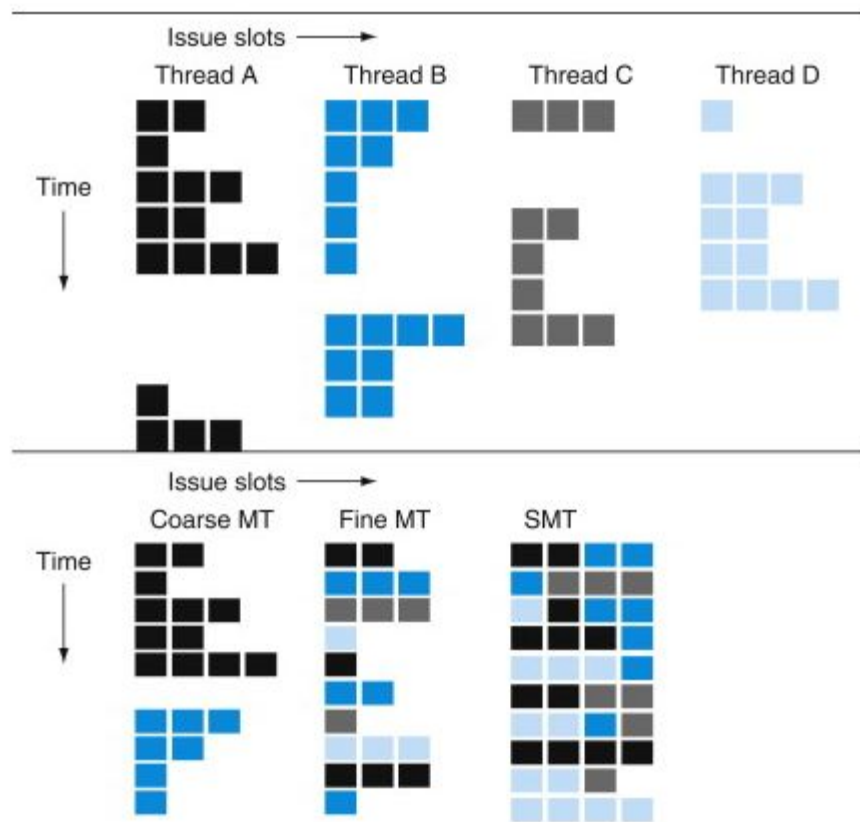
- Operating System further decomposes processes into threads by fragmenting the stack
- Each thread has its own piece of the stack, but they share text, data, and heap
- Threads that share this in a process allow for very cheap context switching



Threads:

- Threads are "light weight processes"
- Threads share Process state among multiple threads ... this greatly reduces the cost of switching context.



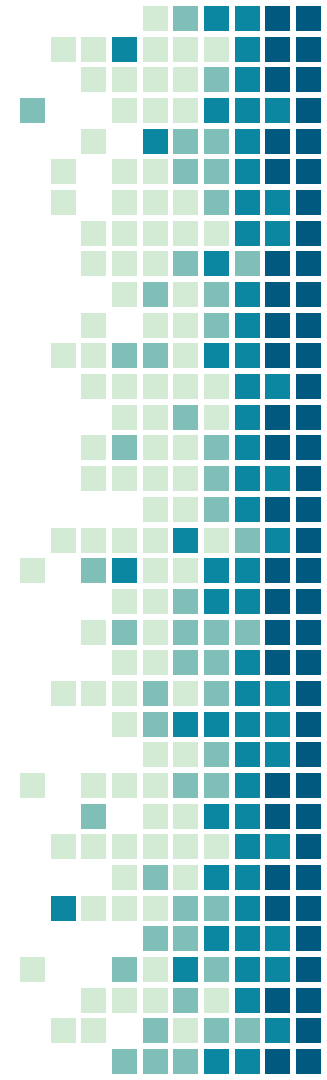


Multithreading in OpenMP

- Single process can spawn multiple threads
- Can have more threads than processor cores
- Threads swap in and out, interleaved in time in all possible ways because of dynamic scheduling capability of the OS
- Have to check program for all possible thread interleaves!



```
[isaac@X550ZA Parallel]$ ./a.out
Hello from thread 3
Hello from thread 0
Hello from thread 2
Hello from thread 1
[isaac@X550ZA Parallel]$ ./a.out
Hello from thread 3
Hello from thread 1
Hello from thread 2
Hello from thread 0
[isaac@X550ZA Parallel]$ ./a.out
Hello from thread 0
Hello from thread 1
Hello from thread 3
Hello from thread 2
[isaac@X550ZA Parallel]$ ./a.out
Hello from thread 2
Hello from thread 0
Hello from thread 3
Hello from thread 1
[isaac@X550ZA Parallel]$ ./a.out
Hello from thread 3
Hello from thread 0
Hello from thread 1
Hello from thread 2
[isaac@X550ZA Parallel]$ ./a.out
Hello from thread 3
Hello from thread 0
Hello from thread 2
Hello from thread 1
[isaac@X550ZA Parallel]$ ./a.out
Hello from thread 1
Hello from thread 0
Hello from thread 3
Hello from thread 2
```



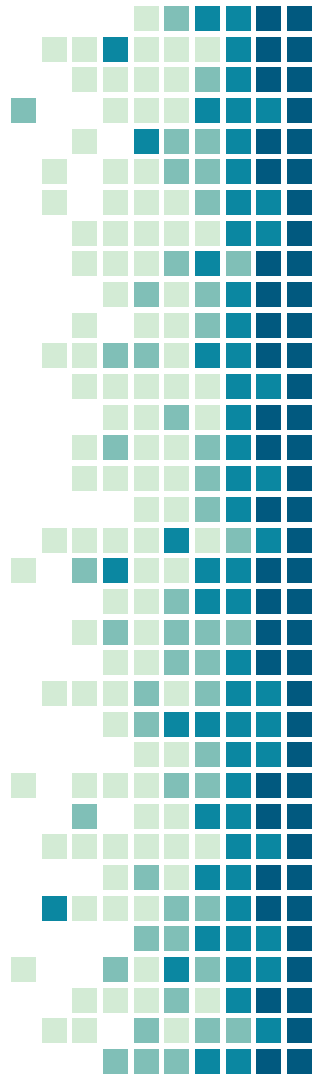
Race Conditions

- Running a program and getting a different answer each time
- Caused by threads sharing data in the heap in a bad way
- Control access to shared variables by **synchronization**




Synchronization

- Protect and order access to shared variables from different threads
- Expensive!
- Synchronise as little as possible by correctly managing the data environment



Fork-Join Parallelism with `#pragma omp parallel`



What the compiler does

```
#pragma omp parallel num_threads(4)
{
    foobar();
}
```

```
void thunk ()
{
    foobar ();
}

pthread_t tid[4];
for (int i = 1; i < 4; ++i)
    pthread_create (
        &tid[i], 0, thunk, 0);
thunk();

for (int i = 1; i < 4; ++i)
    pthread_join (tid[i]);
```

```
#include <stdio.h>
#include <omp.h>

int main()
{
    double array[1000];           // visible to all threads
    omp_set_num_threads(16);      // request 16 threads
    #pragma omp parallel          // fork to 16 threads
    {
        int ID = omp_get_thread_num(); // allocated on the thread's individual stack
                                        // local to the thread, "private"

        printf("Hello from thread %d\n", ID);
    }

    return 0;
}
```

Requesting threads / data management

Sequential code

```
for(i=0;i<N;i++) { a[i] = a[i] + b[i];}
```

OpenMP parallel region

```
#pragma omp parallel
{
    int id, i, Nthrds, istart, iend;
    id = omp_get_thread_num();
    Nthrds = omp_get_num_threads();
    istart = id * N / Nthrds;
    iend = (id+1) * N / Nthrds;
    if (id == Nthrds-1) iend = N;
    for(i=istart;i<iend;i++) { a[i] = a[i] + b[i];}
}
```

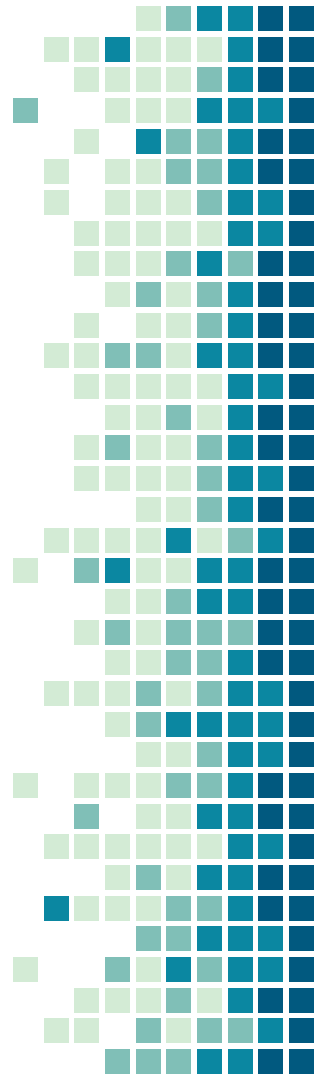
OpenMP parallel region and a worksharing for construct

```
#pragma omp parallel
#pragma omp for
    for(i=0;i<N;i++) { a[i] = a[i] + b[i];}
```

Synchronization

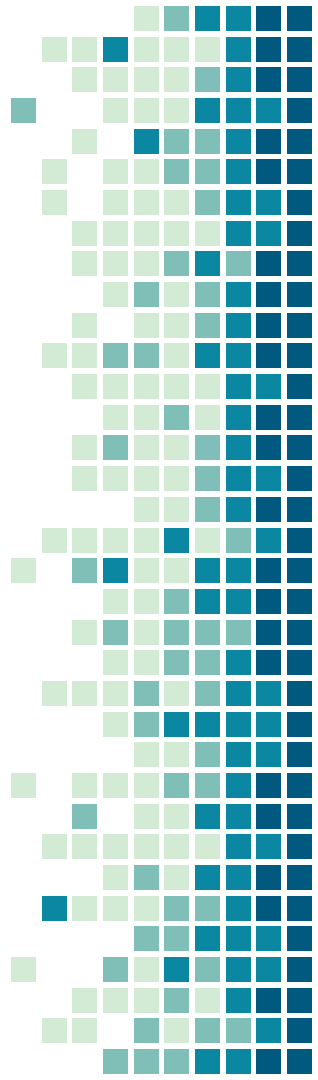
Bringing one or more threads to a well defined and known point in their execution (protect data conflicts, reduce race conditions)

- Barrier: each thread waits at the barrier until all threads arrive
- Mutual exclusion: define a block of code that only one thread at a time can execute



Levels of Synchronization

- High level
 - Critical
 - Atomic
 - Barrier
 - Ordered
- Low level
 - Flush
 - Locks



- Barrier - each thread waits until all the others arrive
 - `#pragma omp barrier`
- Critical - specifies a region that only one thread at a time can enter
 - `#pragma omp critical`
- Atomic - provides mutual exclusion when updating a memory location
 - `#pragma omp atomic`

Additional Resources

- **Excellent book and course on High Performance Computing available entirely for free:**
<https://bitbucket.org/VictorEijkhout/hpc-book-and-course/>
- **OpenMP documentation:** <https://www.openmp.org/>
- **OpenMP tutorial:** <https://computing.llnl.gov/tutorials/openMP/>
- **StackOverflow discussion on concurrency vs parallelism:**
<https://stackoverflow.com/questions/1050222/what-is-the-difference-between-concurrency-and-parallelism>